

✓ 1. They're Redundant (Derived Columns Cover Them)

We engineered two new features:

- `orig_diff = oldbalanceOrig - newbalanceOrig`
- `dest_diff = newbalanceDest - oldbalanceDest`

These capture the **actual transfer behavior**:

- `orig_diff` tells us **how much money left** the sender's account.
- `dest_diff` tells us **how much money entered** the receiver's account.

Once these are created and properly transformed (log1p, clipped), the original columns are **no longer necessary** because:

- They **don't add new information**
 - They could **confuse the model** or add noise due to multicollinearity
-

✓ 2. Some Original Columns Can Be Noisy or Misleading

For example:

- A balance might be zero **before or after** due to temporary internal processing — that doesn't always mean fraud.
- `oldbalanceOrig` and `newbalanceOrig` often contain 0 for system-generated "CASH_IN" or bot-driven transactions.

Using engineered deltas (differences) focuses the model on **actual movement** of funds rather than raw values, which can vary widely.

✓ 3. Model Simplicity & Reduced Dimensionality

Including all original and derived features would result in:

- More input dimensions (risk of overfitting)
- Possible **correlated features** (bad for linear models like Logistic Regression)

By using `amount_log`, `orig_diff_log`, and `dest_diff_log`, we:

- Normalize the skewed distributions
- Keep the most **predictive and interpretable** columns
- Improve generalization

Summary Table

Column	Kept?	Why?
<code>amount</code>	✓	Key feature for fraud, log-transformed
<code>oldbalanceOrig</code>	✗	Replaced by <code>orig_diff</code>
<code>newbalanceOrig</code>	✗	Replaced by <code>orig_diff</code>
<code>oldbalanceDest</code>	✗	Replaced by <code>dest_diff</code>
<code>newbalanceDest</code>	✗	Replaced by <code>dest_diff</code>
<code>orig_diff (new)</code>	✓	Captures sender behavior
<code>dest_diff (new)</code>	✓	Captures receiver behavior
<code>type</code>	✓	Important categorical feature

 **So what now?**

- **If you're using tree-based models** (like Random Forest or XGBoost):
 - **No need to scale** at all! Trees don't care about feature scale.
 - But we often keep scaling in a pipeline for compatibility across models.
- **If you're using Logistic Regression or SVM:**
 - Scaling **is important**, because these models are sensitive to the scale of features.
 - Since `log1p()` helps, but doesn't perfectly normalize, you can:
 - Use `StandardScaler` if your `log1p`-transformed features look roughly normal (check histograms!)
 - Use `MinMaxScaler` if you want everything **strictly in [0,1]**, e.g. for neural nets or interpretability